

Case Study: Finding Real Concurrency Bugs with LLP Predicate Detection

Table of Contents

1. [Overview](#)
2. [Case Study 1: Data Races in Boruvka's Minimum Spanning Tree](#)
 - [2.1 The Algorithm](#)
 - [2.2 The Bug](#)
 - [2.3 Detection Results](#)
 - [2.4 Root Cause Analysis](#)
 - [2.5 The Fix](#)
3. [Case Study 2: Buffer Overflow in a Concurrent Writer](#)
 - [3.1 The Program](#)
 - [3.2 The Bug: TOCTOU Race Leading to Overflow](#)
 - [3.3 Detection Results](#)
 - [3.4 Connecting to RIP Overwrite](#)
 - [3.5 The Fix](#)
4. [Case Study 3: Lost Updates in a Racy Counter](#)
 - [4.1 The Program](#)
 - [4.2 Detection Results](#)
 - [4.3 Quantifying the Damage](#)
 - [4.4 The Fix](#)
5. [Case Study 4: Producer-Consumer Queue Without Synchronization](#)
 - [5.1 The Program](#)
 - [5.2 Detection Results](#)
 - [5.3 The Cascading Failure](#)
 - [5.4 The Fix](#)
6. [Comparison: What Traditional Testing Misses](#)
7. [Reproducing These Results](#)

1. Overview

This document presents four case studies demonstrating the LLP Predicate Detection System's ability to find concurrency bugs that are difficult or impossible to detect through conventional testing. Each case study walks through a real program, explains the bug, shows the detector's output, and analyzes the root cause.

The programs used in these case studies are not contrived examples. The Boruvka MST implementation is a real parallel algorithm from a graduate course project. The buffer overflow, racy counter, and producer-consumer examples are simplified versions of patterns that appear regularly in production concurrent code - shared mutable arrays, unsynchronized counters, and lock-free queues.

What makes these bugs particularly dangerous is their non-determinism. Running the same program 100 times might produce correct results 98 times and silently corrupt data twice. The LLP detector finds these

bugs from a single execution by reasoning about all possible thread interleavings consistent with the observed event ordering.

2. Case Study 1: Data Races in Boruvka's Minimum Spanning Tree

2.1 The Algorithm

Boruvka's algorithm finds the minimum spanning tree (MST) of a weighted graph. It works iteratively: in each round, every connected component finds its cheapest outgoing edge, and these edges are added to the MST. The algorithm terminates when all vertices are in a single component.

The implementation at `parallelalgorithms.group9.homework3.llp.boruvka.Boruvka` uses Java's `parallelStream()` to process edges in parallel during the "find cheapest edge" phase. Each thread reads the `parent[]` array (union-find structure) to determine which component each vertex belongs to, then updates the `cheapest[]` array with the index of the minimum-weight edge for that component.


2.2 The Bug

The `parent[]` and `cheapest[]` arrays are plain `int[]` arrays, accessed from the ForkJoinPool's worker threads through `parallelStream().forEach()` with no synchronization. This creates two classes of data races:

1. **Write-write race on `cheapest[]`:** Two threads processing different edges may both try to update `cheapest[u]` for the same component representative `u`. One thread's update is silently lost, causing the algorithm to select a non-minimum edge.
2. **Read-write race on `parent[]`:** The `find()` method reads `parent[]` to traverse the union-find tree while `union()` writes to it. Path compression during `find()` also writes to `parent[]`, creating a read-write race that can corrupt the tree structure.

2.3 Detection Results

Running the detector against the Boruvka edge-processing phase produces the following output:



The detector found 2 data races on the `cheapest[]` array. The timeline shows the ForkJoinPool worker threads reading `parent[]` and writing `cheapest[]` concurrently, with red **RACE** markers on every event. The contention heat map reveals that `cheapest[0]` and `cheapest[1]` are accessed by multiple worker threads with concurrent vector clocks.

Critically, the detector identifies the specific sequence numbers of the racing events and the exact thread names involved. The execution trace file records the complete scheduling path, enabling a developer to

understand not just *that* a race exists, but *how* the threads interleaved to produce it.

2.4 Root Cause Analysis

The root cause is a fundamental misuse of `parallelStream()`. Java's parallel streams divide work across the ForkJoinPool's worker threads, which execute concurrently. When the lambda passed to `forEach()` accesses shared mutable state (the `cheapest[]` and `parent[]` arrays), every access is a potential data race.

The issue is subtle because `parallelStream()` looks syntactically identical to `stream()`, and the code works correctly with sequential streams. The bug only manifests when the ForkJoin pool assigns the same array element to multiple worker threads -- which depends on the graph structure, the number of available processors, and the pool's work-stealing behavior. This is exactly the kind of bug that passes unit tests on a developer's laptop but causes silent data corruption in production on a 32-core server.

2.5 The Fix

There are two correct approaches:

Option 1: Atomic arrays. Replace `int[] cheapest` with `AtomicIntegerArray` and use `compareAndSet()` to atomically update only if the current value is worse than the proposed value.

```
AtomicIntegerArray cheapest = new AtomicIntegerArray(numNodes);  
// In the parallel forEach:  
cheapest.accumulateAndGet(u, edgeIndex, (current, proposed) ->  
    edges.get(proposed).weight < edges.get(current).weight ? proposed :  
    current);
```

Option 2: Sequential processing. Use `stream()` instead of `parallelStream()` if the graph is small enough that parallelism doesn't provide a meaningful speedup. For graphs with fewer than ~10,000 edges, the overhead of the ForkJoinPool often exceeds the benefit of parallelism.

3. Case Study 2: Buffer Overflow in a Concurrent Writer

3.1 The Program


The `BufferOverflowDemo` program simulates a common pattern in systems programming: multiple threads writing to a shared buffer using a shared write index. The buffer has a fixed size of 4 elements, and two writer threads each attempt to write 3 elements, for a total of 6 writes to a 4-element buffer.


```
static int[] buffer = new int[4];  
static int writeIndex = 0;  
  
// Each thread:  
int idx = writeIndex;           // Read current index  
buffer[idx] = value;           // Write to buffer  
writeIndex = idx + 1;          // Increment index
```

3.2 The Bug: TOCTOU Race Leading to Overflow

The bug is a classic time-of-check-to-time-of-use (TOCTOU) race. Thread A reads `writeIndex` and gets 2. Thread B also reads `writeIndex` and gets 2 -- before Thread A has incremented it. Both threads write to `buffer[2]` (lost update), then both increment `writeIndex` to 3 (lost increment -- should be 4). The counter is now wrong, and subsequent writes will be placed at incorrect indices.

As the threads continue, the `writeIndex` falls behind the actual number of writes performed. Eventually, a thread reads `writeIndex = 4`, which is still within the bounds check `if (idx < buffer.length)`, but because prior increments were lost, there are more writes remaining than buffer slots available. The thread writes to `buffer[4]` -- one element past the end of the buffer.





3.3 Detection Results

The LLP detector finds 12 violations in 23 milliseconds:

```
[LLP] Captured 24 events across 2 threads
[LLP] Detection completed in 23ms
[LLP] Found 0 race(s), 12 violation(s)
```

Violations:


1. BufferOverflow: Concurrent array access + bounds modification:
Writer-1 accesses `buffer[0]` while Writer-2 modifies `writeIndex`
2. BufferOverflow: Concurrent array access + bounds modification:
Writer-1 accesses `buffer[4]` while Writer-2 modifies `count`
- ...

The detector identifies both the direct out-of-bounds access (`buffer[4]` when size is 4) and the concurrent array+index race that caused it. The contention heat map shows `writeIndex` accessed by both threads (6 accesses each) with **SHARED** status, confirming that the index variable is the root cause of the contention.

3.4 Connecting to RIP Overwrite

In native code (C/C++), writing past the end of a stack-allocated buffer overwrites the saved frame pointer and return address. The LLP detector flags this pattern because the combination of an out-of-bounds array write and a concurrent index modification is precisely the mechanism that produces stack-based buffer overflow exploits.

Buffer Overflow Leading to RIP Overwrite



In the diagram above, `buffer[4]` overwrites the saved RBP (frame pointer) and `buffer[5]` would overwrite the return instruction pointer (RIP). While Java's runtime prevents actual memory corruption

through `ArrayIndexOutOfBoundsException`, the detector identifies programs where the *logic* would produce an overflow in a native context. This is valuable for:

- Java code that interfaces with native libraries via JNI or JNA
- Java programs that are prototypes for eventual C/C++ implementations
- Security auditing of concurrent buffer management code

3.5 The Fix

The fix requires atomically reading the index, checking bounds, writing the element, and incrementing the index as a single indivisible operation:

```
synchronized (lock) {
    if (writeIndex < buffer.length) {
        buffer[writeIndex] = value;
        writeIndex++;
    }
}
```

Alternatively, use an `AtomicInteger` for the index and `compareAndSet()` to ensure that only one thread claims each slot:

```
int idx;
do {
    idx = writeIndex.get();
    if (idx >= buffer.length) return; // buffer full
} while (!writeIndex.compareAndSet(idx, idx + 1));
buffer[idx] = value;
```

4. Case Study 3: Lost Updates in a Racy Counter

4.1 The Program

The `RacyCounter` program increments a shared integer counter from 4 threads, each performing 10 increments. The expected final value is 40, but the actual value is consistently lower due to lost updates.

```
static int counter = 0;
// Each thread, 10 times:
int val = counter;           // Read
counter = val + 1;           // Write (based on stale read)
```

4.2 Detection Results

The detector captures 160 events across 4 threads and finds 2,244 data races. The contention heat map is striking:

Location	Counter-0	Counter-1	Counter-2	Counter-3	Status
RacyCounter.counter	 20	 20	 20	 20	RACE(4t)
RacyCounter.buckets[0]	 6	 4	 4	 6	RACE(4t)
RacyCounter.buckets[1]	 6	 6	 4	 4	RACE(4t)

Every field is accessed by all 4 threads with concurrent vector clocks. The `counter` field alone accounts for 80 events (20 per thread: 10 reads + 10 writes), and every inter-thread pair of accesses is a race.

4.3 Quantifying the Damage

The program consistently reports 15-25 lost updates out of 40 expected. The variability itself is diagnostic: a deterministic bug would produce the same wrong answer every time. The fact that the final counter varies between runs confirms that the result depends on thread scheduling, which is the hallmark of a data race.

The execution trace captures the exact scheduling that produced each run's result. By examining the trace, a developer can see where two threads read the same value of `counter`, both computed `val + 1`, and both wrote the same result -- losing one increment entirely.

4.4 The Fix

```
static AtomicInteger counter = new AtomicInteger(0);
// Each thread:
counter.incrementAndGet(); // Atomic read-modify-write
```

5. Case Study 4: Producer-Consumer Queue Without Synchronization

5.1 The Program

The `ProducerConsumer` program implements a circular buffer queue with 2 producer threads and 2 consumer threads. Producers write items to `buffer[tail % size]` and increment `tail` and `count`. Consumers read from `buffer[head % size]` and increment `head` and decrement `count`. None of these operations are synchronized.

5.2 Detection Results

The detector finds 1,082 data races across 4 fields:

Field	Races	Threads
<code>count</code>	412	All 4 threads (producers increment, consumers decrement)
<code>tail</code>	198	Both producers
<code>head</code>	198	Both consumers
<code>buffer[i]</code>	274	Producers write, consumers read same indices

The heat map shows `count` as the most contended field, with all 4 threads reading and writing it concurrently. This is expected: `count` is the shared variable that coordinates production and consumption, and without synchronization, every access is a race.

5.3 The Cascading Failure

What makes this case study particularly instructive is the cascading nature of the bugs. The race on `count` causes producers to overfill the buffer (they check `count < buffer.length` but the stale value of `count` doesn't reflect recent consumer activity). The race on `tail` causes two producers to write to the same slot, losing one item. The race on `head` causes two consumers to read the same slot, processing one item twice.

The detector identifies all three failure modes from a single execution. The execution trace shows the exact scheduling that caused each failure, and the seed value enables reproduction of the same scheduling on the same hardware.

5.4 The Fix

The correct fix depends on performance requirements:

Synchronized queue: Wrap all operations in a single lock. Simple but limits throughput to one operation at a time.

java.util.concurrent.ArrayBlockingQueue: A production-ready bounded queue with separate locks for producers and consumers, plus condition variables for blocking when full or empty.

Lock-free ring buffer: Use `AtomicInteger` for `head` and `tail`, and `compareAndSet()` loops for producers and consumers. Highest throughput but most complex to implement correctly.

6. Comparison: What Traditional Testing Misses

The following table summarizes what each testing approach would find for each case study:

Bug	Unit Test	Stress Test	ThreadSanitizer	LLP Detector
Boruvka data race	Passes (single thread)	May fail intermittently	Detects (C/C++ only)	Detects: 2 races on <code>cheapest[]</code>, <code>parent[]</code>
Buffer overflow TOCTOU	Passes (single thread)	Throws AIOOBE sometimes	Detects (C/C++ only)	Detects: 12 violations, concurrent array+index
Racy counter	Passes (off by 1-20)	Fails (obviously wrong)	Detects (C/C++ only)	Detects: 2,244 races, heat map shows contention
Producer-consumer	Passes (low contention)	May deadlock or lose items	Detects (C/C++ only)	Detects: 1,082 races across 4 fields

The key advantage of the LLP detector over stress testing is *determinism*: stress testing might find a bug after 1,000 runs, or it might not find it at all. The LLP detector finds all races from a single execution by reasoning about the causal structure of events, not by hoping for a lucky thread schedule.

The key advantage over ThreadSanitizer is *platform independence*: ThreadSanitizer is a compile-time instrumentation tool for C/C++ that is not available for Java. The LLP detector works with any Java program on any platform, using pure Java with no native dependencies.

7. Reproducing These Results

All case studies can be reproduced with the following commands:

```
cd my-app

# Case Study 1: Boruvka MST
mvn -q exec:java -Dexec.mainClass=com.utece.student.llpdetection.LLPMain \
  -Dexec.args="com.utece.student.llpdetection.examples.RacyCounter --cli --predicates race"

# Case Study 2: Buffer Overflow
mvn -q exec:java -Dexec.mainClass=com.utece.student.llpdetection.LLPMain \
  -Dexec.args="com.utece.student.llpdetection.examples.BufferOverflowDemo --gui --predicates buffer"

# Case Study 3: Racy Counter
mvn -q exec:java -Dexec.mainClass=com.utece.student.llpdetection.LLPMain \
  -Dexec.args="com.utece.student.llpdetection.examples.RacyCounter --gui --predicates race"

# Case Study 4: Producer-Consumer
mvn -q exec:java -Dexec.mainClass=com.utece.student.llpdetection.LLPMain \
  -Dexec.args="com.utece.student.llpdetection.examples.ProducerConsumer --cli --predicates race"
```

The `--gui` flag generates an interactive HTML report at `llp-report.html`. When violations are found, execution traces are saved as `llp-trace-*.txt` (human-readable) and `llp-trace-*.ser` (serialized, reloadable).

Due to the non-deterministic nature of thread scheduling, the exact number of races and the specific thread interleavings will vary between runs. However, the *types* of violations and the *locations* of the races will be consistent across all runs, because the bugs are structural -- they exist in every execution, even when they don't cause visible failures.